

In-Context Learning: Mechanisms, Applications, and Recent Advances

Generated by AutoSurvey

Contents

1	Introduction to In-Context Learning	2
1.1	Definition and Key Characteristics of In-Context Learning	2
1.2	Emergence of In-Context Learning in Large Language Models	2
1.3	Practical Applications of In-Context Learning	3
1.4	Connections to Other Learning Paradigms	5
2	Theoretical Foundations of In-Context Learning	6
2.1	Inductive Biases and Representation Learning in In-Context Learning	6
2.2	Optimization and Algorithmic Perspectives on In-Context Learning	7
2.3	Connections to Cognitive Science and Neuroscience	8
3	Applications and Empirical Findings of In-Context Learning	9
3.1	Multimodal In-Context Learning for Vision Tasks	9
3.2	Leveraging Text-Image Prompts for In-Context Learning	10
3.3	Benchmarking and Evaluating Multimodal In-Context Learning	11
4	Advancing In-Context Learning Capabilities	11
4.1	Prompt Optimization	11
4.2	Example Selection	12
4.3	Meta-Learning Techniques	14
4.4	Addressing Label Biases	15
5	Practical Considerations and Future Directions	16
5.1	Practical Implications of In-Context Learning	16
5.2	In-Context Learning and Model Architecture	17
5.3	Overcoming Learning Plateaus in In-Context Learning	18
5.4	Theoretical Foundations of In-Context Learnability	19
	References	20

1 Introduction to In-Context Learning

1.1 Definition and Key Characteristics of In-Context Learning

In-context learning is a novel learning paradigm that has emerged alongside the development of large language models (LLMs) [1, 2]. It refers to the remarkable ability of these models to adapt to new tasks and perform them effectively by conditioning on a few demonstration examples, without requiring any updates to their model parameters [3].

At the core of in-context learning is the capacity of a language model to leverage a prompt sequence, consisting of a few input-output pairs (demonstration examples), along with a new query input, to directly generate the corresponding output [4]. This is in contrast to the traditional approach of fine-tuning or updating the model’s parameters on a specific task, which is often time-consuming and computationally intensive.

The key distinguishing feature of in-context learning is its ability to rapidly adapt to new tasks using only a handful of demonstration examples, without the need for any parameter updates [5]. This adaptability is a remarkable feature, as it allows language models to exhibit a high degree of versatility and flexibility, capable of tackling a wide range of tasks and domains with minimal additional training [6].

The emergence of in-context learning capabilities in large language models is a direct consequence of the recent advancements in model scaling, multi-task training, and architectural innovations [7]. As models have grown increasingly larger and more complex, they have developed the ability to capture deeper semantic and reasoning capabilities, which can then be effectively leveraged for in-context learning [3].

A crucial aspect of in-context learning is its reliance on the model’s internal representations and inductive biases, which are shaped during the pretraining phase [2]. These learned representations and biases enable the model to effectively leverage the information contained in the demonstration examples, allowing it to rapidly acquire new skills and apply them to novel inputs [8].

This ability to adapt to new tasks without requiring parameter updates has significant practical implications, as it allows for more efficient and flexible deployment of language models in real-world applications [9]. Instead of the time-consuming and resource-intensive process of fine-tuning or retraining the model for each new task, in-context learning enables a more seamless and on-demand adaptation, potentially opening up new avenues for the widespread adoption of these powerful language models [10].

However, the emergence of in-context learning capabilities is not without its challenges. Researchers have observed that the effectiveness of in-context learning can be highly sensitive to the choice and quality of the demonstration examples provided in the prompt [11]. Factors such as the diversity, relevance, and distribution of the examples can have a significant impact on the model’s ability to learn and generalize effectively [12].

To address these challenges, recent research has explored various strategies, such as prompt optimization, example selection, and meta-learning techniques, to improve the robustness and effectiveness of in-context learning [13]. These advancements aim to make in-context learning more reliable, reproducible, and less sensitive to the specific prompt configurations [14].

1.2 Emergence of In-Context Learning in Large Language Models

The emergence of in-context learning capabilities in large language models (LLMs) can be attributed to several key advancements in the field of natural language processing (NLP) over the past decade. The ability of LLMs to adapt to new tasks through a few demonstration examples, without updating the model’s parameters, has been a remarkable and unexpected development [15].

The historical context leading to the emergence of in-context learning can be traced back to the rapid progress in the scale and capabilities of language models. The success of transformers, such as BERT [16] and GPT [17], has shown the potential of large-scale pre-training on vast amounts of text data to capture rich linguistic knowledge and enable strong performance on a wide range of NLP tasks.

A key driver of the emergence of in-context learning has been the scaling of language models, both in terms of model size and the quantity of training data. As LLMs grew in size, with some models reaching billions of parameters [18], their ability to encode and leverage background knowledge in novel tasks has expanded significantly. The increased capacity of these models has enabled them to learn complex task-agnostic representations that can be quickly adapted to new scenarios through a few demonstration examples [19].

In addition to model scaling, advancements in multi-task training have also played a crucial role in the emergence of in-context learning. By training LLMs on a diverse set of tasks and datasets during pretraining [20], these models have developed a more versatile and generalizable understanding of language, which can then be leveraged for rapid adaptation to new tasks through in-context learning [21].

The architectural innovations in LLMs, particularly the self-attention mechanism in transformers, have also been instrumental in enabling in-context learning. The self-attention mechanism allows the model to dynamically attend to relevant parts of the input, including the demonstration examples, and seamlessly integrate this information with its existing knowledge to generate appropriate outputs for the new task [13].

Furthermore, the development of techniques for prompt engineering and prompt tuning [3] has further enhanced the ability of LLMs to effectively utilize the contextual information provided in the form of task demonstrations. By optimizing the prompts, researchers have been able to unlock the in-context learning capabilities of these models and push the boundaries of what they can achieve without fine-tuning [22].

The emergence of in-context learning in LLMs has also been linked to the models' ability to perform implicit Bayesian inference and uncover the latent structure underlying the provided examples [23]. By leveraging the coherence and compositional nature of language data during pretraining, LLMs can develop the capacity to infer the shared conceptual structure across the demonstration examples and rapidly apply this knowledge to the new task at hand.

Overall, the emergence of in-context learning in large language models has been driven by a confluence of factors, including the scaling of model size and data, advancements in multi-task training, architectural innovations, and the development of techniques for prompt engineering. These developments have enabled LLMs to exhibit remarkable adaptability and versatility, paving the way for a new paradigm of learning that can have far-reaching implications across various application domains [24].

1.3 Practical Applications of In-Context Learning

The rapid advancements in large language models (LLMs) have enabled the emergence of in-context learning (ICL) as a powerful paradigm with a wide range of practical applications across diverse domains [25]. This capability has emerged as a result of several key developments in the field of natural language processing (NLP) over the past decade, including the scaling of language models, advancements in multi-task training, and architectural innovations.

The scaling of language models, both in terms of model size and the quantity of training data, has been a crucial driver of the emergence of in-context learning [18]. As LLMs grew in size, with some models reaching billions of parameters, their ability to encode and leverage background

knowledge in novel tasks has expanded significantly. The increased capacity of these models has enabled them to learn complex task-agnostic representations that can be quickly adapted to new scenarios through a few demonstration examples [19].

In addition to model scaling, advancements in multi-task training have also played a crucial role in the emergence of in-context learning. By training LLMs on a diverse set of tasks and datasets during pretraining [20], these models have developed a more versatile and generalizable understanding of language, which can then be leveraged for rapid adaptation to new tasks through in-context learning [21].

The architectural innovations in LLMs, particularly the self-attention mechanism in transformers, have also been instrumental in enabling in-context learning. The self-attention mechanism allows the model to dynamically attend to relevant parts of the input, including the demonstration examples, and seamlessly integrate this information with its existing knowledge to generate appropriate outputs for the new task [13].

Furthermore, the development of techniques for prompt engineering and prompt tuning [3] has further enhanced the ability of LLMs to effectively utilize the contextual information provided in the form of task demonstrations. By optimizing the prompts, researchers have been able to unlock the in-context learning capabilities of these models and push the boundaries of what they can achieve without fine-tuning [22].

ICL has found applications in a wide range of domains, including natural language processing, computer vision, and scientific machine learning. In the field of NLP, LLMs have demonstrated the ability to perform various language tasks, such as text generation, question answering, and sentiment analysis, by simply conditioning on a few relevant examples in the input [26]. This capability has significant practical implications, as it allows for rapid task adaptation and enables the deployment of these models in dynamic, real-world scenarios where the task requirements may change frequently.

Beyond NLP, ICL has also found applications in the computer vision domain, where vision-language models (VLMs) can leverage multimodal in-context examples to perform tasks like image captioning, object detection, and semantic segmentation [27]. The ability to adapt to new visual tasks with only a few demonstrations can be particularly valuable in scenarios where data collection and labeling are costly, or when the task requirements evolve over time.

The versatility of ICL is further exemplified by its application in the domain of egocentric video understanding, as explored in the paper “Efficient In-Context Learning in Vision-Language Models for Egocentric Videos”. The authors propose a training method called “Efficient In-context Learning on Egocentric Videos (EILEV)”, which enables VLMs to effectively learn in-context on egocentric video data, even with limited training resources. This approach demonstrates the potential of ICL to adapt to specialized domains and data modalities, where traditional fine-tuning methods may struggle due to data scarcity or domain shifts.

Moreover, the paper “Skeleton-in-Context: Unified Skeleton Sequence Modeling with In-Context Learning” investigates the application of ICL to the task of skeleton sequence modeling, which is crucial for various applications in computer vision and robotics. The authors propose a novel framework, Skeleton-in-Context (SiC), that can handle multiple skeleton-based tasks simultaneously and generalize to new, unseen tasks by leveraging in-context learning. This work highlights the versatility of ICL and its ability to be applied to diverse data modalities beyond just text and images.

While the broad adoption of ICL holds great promise, the successful application of this paradigm also comes with its own set of challenges and considerations. The performance of ICL-based models can be highly sensitive to the quality and relevance of the in-context examples provided, as well as the design of the prompts used to guide the learning process [28]. Addressing these challenges, through techniques such as prompt optimization, example selection, and meta-learning, is an active

area of research that can help unlock the full potential of ICL in practical applications.

Overall, the diverse applications of ICL across natural language processing, computer vision, and scientific machine learning domains demonstrate the versatility and promise of this paradigm. As the field continues to evolve, the ability of ICL to enable rapid task adaptation and efficient knowledge transfer can have far-reaching implications for the development of more adaptable and versatile AI systems that can thrive in complex, dynamic environments.

1.4 Connections to Other Learning Paradigms

In-context learning has strong connections to several other prominent learning paradigms, including meta-learning, few-shot learning, and Bayesian inference. These paradigms share underlying principles and techniques that can enable cross-pollination of ideas and further advancements in the field of in-context learning.

Meta-learning, or learning to learn, is a closely related framework that aims to develop models that can rapidly adapt to new tasks using a limited amount of training data [29]. In meta-learning, a model is trained on a diverse set of tasks, enabling it to learn general strategies and representations that can be quickly fine-tuned or adapted to novel tasks. This is analogous to the in-context learning ability of large language models, where the models leverage their broad pretraining to quickly learn new tasks from just a few examples provided in the input context.

The connections between in-context learning and meta-learning are further highlighted by the role of inductive biases and representation learning. Both in-context learning and meta-learning rely on the model’s internal representations and inductive biases to enable rapid adaptation to new tasks [30]. The architectural choices, pretraining, and learning dynamics of large language models shape these inductive biases, allowing the models to effectively leverage context information to perform in-context learning. Similarly, meta-learning approaches aim to learn representations and learning algorithms that can be efficiently adapted to new tasks, drawing on related principles.

Another closely related paradigm is few-shot learning, which focuses on the ability to learn new concepts from a small number of examples. In-context learning can be seen as a specific form of few-shot learning, where the model is given a few demonstration examples within the input context to adapt to a new task [31]. The shared goal of rapidly learning from limited data has led to the exploration of techniques that can enable both in-context learning and few-shot learning, such as prompt optimization, example selection, and meta-learning approaches.

Bayesian inference offers a theoretical framework that can provide insights into the mechanisms underlying in-context learning. The ability to rapidly update beliefs and make decisions based on a small amount of data is a hallmark of Bayesian reasoning [32]. In the context of in-context learning, Bayesian models can provide a principled way to incorporate the information from the few provided examples and the model’s prior knowledge to make predictions. Furthermore, the connections between in-context learning and Bayesian inference can be explored through the lens of cognitive science and neuroscience, as human learning often exhibits Bayesian-like characteristics, such as analogical reasoning and rapid concept acquisition [29].

The shared underlying principles and techniques across these learning paradigms suggest that there is significant potential for cross-pollination of ideas and mutual advancement. For example, the Bayesian perspective on in-context learning can inform the development of more robust and uncertainty-aware in-context learning models, while the meta-learning framework can be leveraged to learn more effective in-context learning strategies. Similarly, the insights gained from in-context learning, such as the importance of prompt design and example selection, can be applied to enhance few-shot learning and meta-learning approaches.

Moreover, the connections between in-context learning and other learning paradigms extend

beyond the algorithmic and theoretical aspects. There are also intriguing parallels in the practical applications and the unique challenges faced by these learning approaches. For instance, the ability to effectively leverage multimodal information, such as textual prompts and visual examples, is crucial for in-context learning in computer vision tasks, as well as for few-shot and meta-learning in multimodal settings [33]. Addressing the sensitivity of in-context learning to prompt configurations and label biases is another shared challenge that can benefit from cross-pollination of ideas with meta-learning and Bayesian inference techniques [34].

In conclusion, the connections between in-context learning and other prominent learning paradigms, such as meta-learning, few-shot learning, and Bayesian inference, highlight the potential for synergistic advancements in the field of adaptable and versatile AI systems. By leveraging the shared underlying principles and techniques, researchers can drive progress in in-context learning and further bridge the gap between the remarkable learning capabilities of humans and the current limitations of machine learning models.

2 Theoretical Foundations of In-Context Learning

2.1 Inductive Biases and Representation Learning in In-Context Learning

The emergence of in-context learning (ICL) capabilities in large language models (LLMs) has been driven by a complex interplay between model architecture, pretraining, and the learning dynamics. Examining the inductive biases that enable ICL is crucial to understanding how these models can effectively adapt to new tasks using only a few demonstration examples.

A key factor underlying the ICL abilities of LLMs is the development of internal representations that are conducive to rapid task adaptation. The model architecture, particularly the attention mechanism, plays a critical role in shaping these representations [35]. Transformer-based models, which have become ubiquitous in the field of language modeling, are particularly well-suited for ICL due to their ability to effectively capture complex dependencies and patterns in the input data.

During pretraining on large and diverse datasets, the model’s internal representations are progressively shaped to capture the inductive biases that enable effective in-context learning. This process involves learning robust feature representations, understanding the underlying structure and relationships in the data, and developing the capacity to rapidly integrate new information [36].

One of the critical inductive biases that emerges in LLMs during pretraining is the ability to identify and leverage the relevant features and contextual cues that are predictive of the target output, even in the presence of underspecified or ambiguous demonstration examples [37]. This feature bias allows the model to effectively generalize from a small number of in-context examples, focusing on the most informative aspects of the input and disregarding less relevant or potentially misleading information.

Moreover, the learning dynamics during pretraining play a crucial role in the development of ICL capabilities. Studies have shown that the emergence of ICL is often accompanied by the abrupt formation of an “induction head” within the model’s internal representations, which competes with traditional in-weights learning [38]. This induction head is believed to be responsible for the model’s ability to rapidly adapt to new tasks by leveraging the information contained in the in-context examples.

Interestingly, the inductive biases that enable ICL can be further shaped and enhanced through targeted pretraining strategies. For example, [39] demonstrates that the model’s in-context learning performance can be improved by exposing it to a carefully curated curriculum of examples during pretraining, mirroring the learning dynamics observed in human cognition.

Furthermore, the theoretical connections between ICL and other learning paradigms, such as meta-learning, few-shot learning, and Bayesian inference, provide valuable insights into the underlying mechanisms [40, 41]. These connections suggest that the inductive biases that enable ICL may be related to the model’s ability to effectively represent and reason about probabilistic relationships, and to rapidly adapt its internal representations to new tasks and domains.

Overall, the inductive biases and representation learning processes that enable in-context learning in large language models are multifaceted and deeply intertwined with the model architecture, pretraining, and learning dynamics. By understanding these underlying mechanisms, researchers can work towards developing more versatile and adaptable AI systems that can effectively leverage in-context learning to tackle a wide range of tasks and challenges.

2.2 Optimization and Algorithmic Perspectives on In-Context Learning

The optimization and algorithmic perspectives on in-context learning (ICL) have been actively explored, providing valuable insights into the underlying mechanisms and connections to other learning paradigms. These investigations have delved into gradient-based methods, task-agnostic and task-aware approaches, as well as the links to classical learning algorithms such as Bayesian inference and multi-task learning.

Gradient-based methods have been a focal point in understanding the optimization dynamics of ICL. [42] analyzes the connection between optimization algorithms and physical systems, offering a unified framework to study gradient descent and its variants. This work highlights the potential for drawing parallels between the optimization of ICL models and the dynamics of physical systems, which could lead to a deeper understanding of the underlying mechanisms.

The challenge of conflicting gradients across tasks in multi-task learning has also been addressed. [43] introduces the Conflict-Averse Gradient Descent (CAGrad) algorithm, which aims to minimize the average loss function while leveraging the worst local improvement of individual tasks to regularize the algorithm trajectory. This approach provably converges to a minimum over the average loss, demonstrating its effectiveness on challenging multi-task supervised learning and reinforcement learning tasks.

Beyond gradient-based methods, task-agnostic and task-aware approaches have been investigated. [44] presents a variational Bayesian framework for multi-task learning, where task relatedness is explored by specifying learnable priors on the task-specific parameters. This principled approach enables the model to jointly infer the posteriors of representations and classifiers, exploiting task relatedness in a robust manner.

The connections between ICL and classical learning algorithms, such as Bayesian inference and multi-task learning, have also been studied. [45] introduces a gradient aggregation approach that leverages Bayesian inference to quantify the uncertainty in each gradient dimension, leading to state-of-the-art performance on various datasets.

Furthermore, the relationship between ICL and the learnability of tasks has been explored. [46] establishes a statistical task complexity bound for attention model pretraining, showing that effective pretraining only requires a small number of independent tasks and that the pretrained model closely matches the Bayes optimal algorithm.

The optimization and algorithmic perspectives on ICL have also been studied in the context of reinforcement learning. [47] introduces the In-context Exploration-Exploitation (ICEE) algorithm, which performs an exploration-exploitation trade-off at inference time within a Transformer model, without the need for explicit Bayesian inference.

Overall, the exploration of optimization and algorithmic perspectives on in-context learning has revealed valuable insights into the mechanisms underlying this powerful learning paradigm. These

investigations have uncovered connections to gradient-based methods, task-agnostic and task-aware approaches, as well as classical learning algorithms, paving the way for the development of more effective and robust in-context learning techniques across diverse domains.

2.3 Connections to Cognitive Science and Neuroscience

The remarkable ability of in-context learning exhibited by large language models (LLMs) has drawn intriguing parallels to certain cognitive processes observed in human learning. Exploring these connections can provide valuable insights into the underlying mechanisms of in-context learning and potentially inform the development of more versatile and adaptable AI systems.

One of the key cognitive processes that shares remarkable similarities with in-context learning is analogy-making [48]. Analogical reasoning, the ability to draw connections between seemingly unrelated concepts, is a hallmark of human intelligence and a fundamental aspect of concept formation, problem-solving, and knowledge transfer. In the context of in-context learning, the capacity to rapidly acquire and apply new concepts based on a few demonstration examples resonates with the human capacity for analogy-making.

Computational models of analogy, such as the Structure-Mapping Theory [48], offer a framework for understanding the cognitive mechanisms that enable humans to perceive and exploit structural similarities between concepts. These models emphasize the role of relational representations, where the meaning of a concept is defined not only by its intrinsic features but also by its relationships to other concepts. Similarly, in-context learning in LLMs relies on the ability to extract and leverage relevant relational information from the provided examples, allowing for the rapid acquisition and application of new skills and knowledge.

The human ability to rapidly acquire concepts, often referred to as “fast mapping,” is another cognitive process that bears striking similarities to in-context learning [49]. Fast mapping enables humans to learn the meaning of new words or concepts from minimal exposure, drawing upon their prior knowledge and reasoning capabilities. This capacity for rapid concept acquisition is a hallmark of human cognitive development and is believed to be facilitated by the brain’s ability to form multimodal representations that integrate various sources of information.

Computational models of fast mapping, such as the one proposed in [49], suggest that the coordination of multisensory and text-derived representations, mediated by a semantic control system, is a key mechanism underlying this cognitive ability. Interestingly, the recent advancements in multimodal learning and the integration of textual and visual information in in-context learning [50, 51] echo these principles, highlighting the potential for cross-pollination between research in cognitive science and AI.

Furthermore, the ability to transfer learning from one context to another, a fundamental aspect of human cognition, is closely related to the in-context learning paradigm [52]. Transfer learning, the capacity to leverage knowledge acquired in one domain to improve performance in a related but distinct task, is a crucial cognitive ability that enables flexible and adaptive behavior. Interestingly, the connections between transfer learning and the representation of cognitive task relationships in the brain [52] suggest that the mechanisms underlying in-context learning may share commonalities with the neural mechanisms that support human transfer learning.

Neuroscientific investigations into the cognitive processes underlying in-context learning can also provide valuable insights. Studies exploring the neural correlates of rapid concept acquisition, analogical reasoning, and transfer learning can shed light on the computational principles and representations that enable these abilities [53]. For instance, the role of the hippocampus and associated brain regions in supporting flexible, context-dependent learning and the formation of relational representations [53] may offer clues about the neural mechanisms underlying in-context

learning in LLMs.

Additionally, the interplay between task-specific and task-agnostic representations in the brain [52] may provide inspiration for the design of in-context learning architectures that can effectively leverage both task-specific knowledge and general, transferable representations. By bridging the gap between cognitive science and AI, researchers can uncover the principles that underlie flexible and adaptive learning, ultimately informing the development of more human-like AI systems capable of in-context learning.

3 Applications and Empirical Findings of In-Context Learning

3.1 Multimodal In-Context Learning for Vision Tasks

The rapid advancement of large language models (LLMs) has enabled the emergence of in-context learning (ICL) in the computer vision domain [25, 54]. Unlike traditional approaches that rely on task-specific models and finetuning, ICL allows vision models to adapt to new tasks using only a few demonstration examples without updating the model parameters. This capability has significant implications for developing versatile and adaptable computer vision systems.

One of the key challenges in applying ICL to vision tasks is the inherent multimodal nature of the problem, where both visual and textual information need to be effectively leveraged [55, 56]. Traditional vision tasks, such as semantic segmentation, object detection, and image captioning, often require understanding the semantic and contextual relationships between different visual elements and their linguistic descriptions.

In-context learning for such multimodal vision tasks presents unique opportunities and challenges. On the one hand, the ability to incorporate both visual and textual prompts can potentially enhance the model’s learning and generalization capabilities [55]. By providing relevant example images and task descriptions, the model can more effectively learn to adapt to new instances of the task. On the other hand, the integration of these diverse modalities and the effective utilization of their complementary information pose significant technical hurdles.

Recent research has explored various approaches to address the challenges of multimodal in-context learning for vision tasks. One promising direction is the use of “inpainting-based” prompting, where the model is trained to generate the desired output (e.g., segmentation mask, object bounding box) given the input image and a textual description of the task [55]. By leveraging a generative inpainting mechanism, the model can learn to effectively map the visual and textual information to the target output, enabling versatile in-context adaptation across different vision tasks.

Another approach focuses on unifying the representation of visual and textual information in a shared embedding space, allowing the model to seamlessly handle diverse in-context prompts [25]. By encoding both image and text prompts into a common representational format, the model can perform generative modeling on the interleaved sequence of visual and textual elements, facilitating in-context learning for a wide range of vision-language tasks.

Empirical evaluations of these multimodal in-context learning approaches have demonstrated promising results across various computer vision tasks [55, 25]. For example, the inpainting-based prompting method has shown significant performance improvements in tasks like semantic segmentation and object detection compared to traditional finetuning-based approaches, highlighting the potential of in-context learning to provide efficient and versatile solutions.

3.2 Leveraging Text-Image Prompts for In-Context Learning

In-context learning has emerged as a powerful paradigm that allows language models to adapt to new tasks using only a few demonstration examples, without the need for expensive fine-tuning. While this capability has been extensively studied in the context of natural language processing, recent advancements have extended in-context learning to the multimodal domain, where both textual and visual information can be leveraged to enable effective task adaptation.

One of the key aspects of successful multimodal in-context learning is the ability to effectively leverage both textual prompts and visual examples. Textual prompts can provide high-level task descriptions and instructions, while visual examples can offer concrete demonstrations of the desired outputs. The interplay between these two modalities can lead to synergistic effects, where the combination of text and images can provide more informative in-context cues than either modality alone.

Several recent studies have explored methods for leveraging text-image prompts to enable in-context learning in multimodal settings. For example, the IMProv model [55] proposes a generative transformer-based architecture that can learn to perform a variety of computer vision tasks, such as foreground segmentation and object detection, by conditioning on both textual task descriptions and visual examples. Similarly, the work on “Exploring Effective Factors for Improving Visual In-Context Learning” [56] highlights the importance of prompt selection and prompt fusion for enhancing the performance of in-context learning in computer vision tasks.

The complementary nature of text and image prompts in multimodal in-context learning has also been examined in the context of image captioning. The work “Exploring Diverse In-Context Configurations for Image Captioning” [57] investigates the effects of various strategies for selecting and assigning image-text pairs as in-context examples, demonstrating that the performance of in-context learning for image captioning can be significantly improved by employing effective prompt configuration strategies.

Beyond the specific applications mentioned above, the broader “Understanding and Improving In-Context Learning on Vision-language Models” [28] study provides valuable insights into the role of textual and visual information in multimodal in-context learning. The authors find that the in-context learning performance of vision-language models is predominantly driven by the textual information in the demonstrations, with the visual information having a relatively minor impact. This finding suggests that the effective utilization of textual prompts is crucial for enhancing the in-context learning capabilities of these models, and that the design of multimodal prompts should carefully consider the relative importance of the two modalities.

Building on these insights, the “MMICL: Empowering Vision-language Model with Multi-Modal In-Context Learning” [58] framework proposes a novel approach to enable vision-language models to effectively leverage complex multimodal prompts, including multiple images and text. The authors introduce a context scheme that augments the in-context learning ability of the model, and construct a dataset specifically designed to enhance the model’s understanding of complex multimodal prompts.

Furthermore, the “Lost in Translation: When GPT-4V(ision) Can’t See Eye to Eye with Text. A Vision-Language-Consistency Analysis of VLLMs and Beyond” [59] study takes a closer look at the interplay between the vision and language modalities in large multimodal language models. The authors introduce a systematic framework to quantify the capability disparities between the two modalities and propose a “Vision Description Prompting” method to effectively improve performance on challenging vision-related tasks by leveraging the textual information in the prompts.

In summary, the research on leveraging text-image prompts for in-context learning in multimodal settings highlights the importance of effectively combining textual and visual information to enable

effective task adaptation. The reviewed studies demonstrate that the careful selection, fusion, and configuration of textual prompts and visual examples can lead to significant improvements in the in-context learning performance of vision-language models, across a variety of applications. These findings underscore the necessity for further research into the interplay between language and vision in the context of in-context learning, to unlock the full potential of these models in real-world multimodal tasks.

3.3 Benchmarking and Evaluating Multimodal In-Context Learning

The effective evaluation of multimodal in-context learning capabilities in vision-language models (VLMs) is crucial for driving progress in this rapidly evolving field. While existing benchmarks have made valuable contributions, there is a need for a comprehensive and rigorous evaluation framework that can capture the nuances of multimodal in-context learning.

One prominent benchmark in this space is MMBench [60], which presents a meticulously curated dataset and a novel evaluation strategy to assess the diverse abilities of VLMs. MMBench goes beyond traditional metrics by incorporating the use of ChatGPT to facilitate a more robust evaluation of model predictions. However, MMBench primarily focuses on the overall capabilities of VLMs and does not explicitly target in-context learning.

Another benchmark, CODIS [61], is designed to assess the ability of models to utilize contextual information provided in free-form text to enhance their visual comprehension. CODIS highlights the importance of context-dependent multimodal understanding, which is a crucial aspect of in-context learning. The findings from CODIS suggest that current multimodal models struggle to effectively extract and utilize contextual information, underscoring the need for more advanced in-context learning capabilities.

The IGLUE benchmark [62] offers a comprehensive evaluation of multimodal transfer learning, including zero-shot and few-shot learning settings. While IGLUE does not directly target in-context learning, the few-shot learning setup provides valuable insights into the models’ ability to adapt to new tasks using limited examples, which is a core aspect of in-context learning.

More recently, the MMICL [58] approach has been proposed, which introduces a novel context scheme and a multi-modal in-context learning (MIC) dataset to enhance the in-context learning capabilities of VLMs. The MMICL framework demonstrates state-of-the-art performance on various vision-language tasks, particularly in handling complex multi-modal prompts, highlighting the importance of targeted in-context learning evaluation.

To further advance the benchmarking and evaluation of multimodal in-context learning, we propose the following key considerations: task design, multimodal prompts, evaluation metrics, benchmarking datasets, and comparison to human performance. By adopting these considerations, the research community can develop a robust and comprehensive benchmark for evaluating multimodal in-context learning capabilities, enabling a deeper understanding of the current state-of-the-art and guiding future advancements in this rapidly evolving field.

4 Advancing In-Context Learning Capabilities

4.1 Prompt Optimization

In-context learning (ICL) has shown promising results in adapting large language models (LLMs) to a wide range of tasks with minimal fine-tuning. However, the performance of ICL is highly dependent on the quality of the prompts used to condition the model. Poorly designed prompts

can lead to suboptimal performance, unstable learning, and biased outputs. Therefore, prompt optimization has become a crucial component in advancing the capabilities of in-context learning.

One of the key techniques for prompt optimization is prompt search [63]. Prompt search aims to find the most effective prompt for a given task by exploring the vast space of possible prompts. This can be done through both manual and automated methods. Manual prompt search involves iteratively crafting and evaluating different prompts, leveraging human intuition and domain knowledge. Automated prompt search, on the other hand, employs techniques such as gradient-based optimization [64], evolutionary algorithms, or reinforcement learning [64] to explore the prompt space in a more systematic and scalable manner.

Another approach to prompt optimization is prompt tuning [65, 66]. In contrast to prompt search, which focuses on finding the optimal prompt, prompt tuning aims to learn a task-specific prompt that can be appended to the input during inference. This is typically done by introducing a small number of learnable prompt tokens that are optimized along with the model’s prediction head, while the rest of the model parameters remain frozen. The key advantage of prompt tuning is that it can be applied to a wide range of tasks without the need for manual prompt engineering.

Recent work has also explored prompt encoding [67], which aims to learn a more expressive and informative prompt representation. By encoding the prompt using a neural network, the model can capture complex relationships between the prompt and the task-relevant information, leading to improved in-context learning performance.

One important consideration in prompt optimization is the robustness of the resulting prompts. Prompts that are highly sensitive to small changes in the input or the demonstration examples can lead to unstable and unreliable in-context learning [63]. To address this, researchers have explored methods to improve the robustness of prompts, such as leveraging contrastive learning [68, 69] or incorporating task-relevant information into the prompt [65, 70].

Another important aspect of prompt optimization is the interaction between prompts and in-context examples. Recent studies have shown that the effectiveness of in-context examples can depend on the quality and specificity of the prompt [71]. In some cases, providing detailed task instructions in the prompt can diminish the need for optimizing the in-context examples, as the prompt itself contains sufficient information to guide the model’s learning. Understanding the interplay between prompts and in-context examples is a crucial direction for advancing in-context learning capabilities.

In summary, prompt optimization is a crucial component in improving the effectiveness and robustness of in-context learning. Techniques such as prompt search, prompt tuning, and prompt encoding, along with considerations of prompt robustness and the interaction between prompts and in-context examples, have shown promising results in advancing the capabilities of in-context learning. As the field of in-context learning continues to evolve, further research is needed to develop more sophisticated and generalizable prompt optimization methods that can unlock the full potential of large language models in a wide range of applications.

4.2 Example Selection

In-context learning (ICL) has shown remarkable potential in enabling large language models (LLMs) to adapt to new tasks with just a few demonstration examples. However, the performance of ICL is highly sensitive to the choice of those examples [72, 73]. Selecting the most informative in-context examples is a crucial component in advancing the capabilities of ICL.

One approach to example selection is based on similarity-based retrieval [73]. The key idea is to identify examples that are most similar to the target task or input, under the assumption that such examples will provide the most relevant information to the model. Similarity can be measured using

a variety of metrics, such as BERTScore-Recall (BSR), which has been shown to outperform other standard similarity measures in selecting examples that demonstrate more of the salient aspects of the test input [73].

While similarity-based retrieval can be effective, it may not capture the full breadth of information required for successful ICL. Diversity-based selection approaches have been proposed as a complementary strategy to ensure that the selected examples cover a wider range of relevant information [72, 74]. The underlying principle is that a diverse set of examples can provide the model with a more comprehensive understanding of the task, enabling it to better generalize to new instances.

One method for achieving diversity-based example selection is to formulate the problem as a combinatorial optimization task, where the goal is to find a subset of examples that maximizes a diversity-aware objective function [72]. This can be challenging, as the strong dependencies among the examples make it an NP-hard problem. To address this, researchers have proposed LENS, a two-stage approach that first filters the dataset to obtain informative examples individually, and then performs a diversity-guided search to refine the selected set [72].

Another approach to diversity-based example selection is to model the problem as a Markov decision process and train an example retriever using reinforcement learning [74]. This method, called RetICL, can capture the dependencies between examples and optimize for the sequential selection of examples that maximizes the in-context learning performance. By learning to model the interplay between examples, RetICL can implicitly discover representations of problem-solving strategies that are beneficial for the target task.

In addition to similarity-based and diversity-based methods, example-based optimization techniques have also been explored for in-context example selection [73]. These approaches aim to directly optimize the performance of the in-context learning model by selecting a set of examples that maximizes a specific evaluation metric, such as those that capture both relevance and diversity (e.g., ERR-IA, a-NDCG, and NRBP) [75].

One example of an example-based optimization method is the use of set-level metrics, such as Set-BSR, which extend individual similarity-based metrics to capture the coverage of the salient aspects of the test input by the selected example set [73]. By optimizing for set-level coverage, these methods can outperform independent ranking approaches, even without any task-specific or LLM-specific training.

Beyond these core approaches, researchers have also explored other techniques to further enhance the effectiveness and efficiency of in-context example selection. For instance, the concept of “data curation” has been shown to be a powerful method for stabilizing ICL performance without any other changes to the ICL algorithm [76]. By carefully curating a subset of training data, the variability in ICL performance can be significantly reduced, even when using randomly sampled examples.

Additionally, some studies have investigated the role of prompt optimization and example-based meta-learning in improving in-context learning capabilities [77, 78]. Prompt optimization techniques, such as prompt search, prompt tuning, and prompt encoding, can help to align the in-context examples with the target task, leading to better performance. Meta-learning approaches, on the other hand, can be used to learn better prompt initialization and in-context learning strategies, further enhancing the adaptability of ICL to new tasks.

It is worth noting that the selection of in-context examples is not only important for the performance of ICL but also has implications for the practical deployment of these systems. The trade-offs between performance, efficiency, and interpretability need to be carefully considered, as the choice of examples can impact the reliability, transparency, and overall usability of in-context learning models in real-world applications [13].

In summary, the selection of informative in-context examples is a crucial component in advancing

the capabilities of in-context learning. Researchers have explored a range of approaches, including similarity-based retrieval, diversity-based selection, and example-based optimization, to identify examples that can effectively capture the relevant information for the target task. These methods have shown promising results, but there is still room for further improvement, particularly in addressing the practical challenges of deploying ICL systems in real-world settings.

4.3 Meta-Learning Techniques

Meta-learning approaches have emerged as a promising strategy for addressing the challenges of in-context learning, particularly in the few-shot setting. These approaches aim to leverage prior experience across multiple tasks to learn an effective initialization or adaptation strategy that can be quickly applied to new tasks with limited data.

One key meta-learning technique that has shown promise for in-context learning is few-shot learning. Few-shot learning algorithms, such as Prototypical Networks [79] and MAML (Model-Agnostic Meta-Learning) [80], are designed to learn a model initialization that can be quickly adapted to new tasks with only a few labeled examples. In the context of in-context learning, these few-shot learning methods can be employed to learn a prompt initialization that enables the model to effectively adapt to new tasks through in-context learning.

For example, the MetaICL (Meta-training for In-Context Learning) framework [81] leverages a meta-training approach to optimize the language model’s parameters for effective in-context learning on a diverse set of tasks. By meta-training the model to perform well on a wide range of in-context learning tasks, MetaICL enables the model to quickly adapt to new tasks using only a few demonstration examples, without the need for any parameter updates.

Another meta-learning technique that has shown promise for in-context learning is meta-gradient optimization. Meta-gradient optimization approaches, such as MAML and its variants, aim to learn an effective parameter update rule that can be applied to new tasks to quickly adapt the model. In the context of in-context learning, these meta-gradient optimization methods can be used to learn a prompt update rule that enables the model to efficiently update its in-context representations in response to new demonstration examples.

For instance, the Gradient-RegulATed Meta-prompt learning (GRAM) framework [82] leverages a meta-learning approach to jointly optimize the prompt initialization and a gradient regularization function. This allows the model to quickly adapt its in-context representations to new tasks while maintaining strong cross-domain generalization, addressing the common challenge of prompt tuning techniques overfitting to the limited training data.

Furthermore, recent work has explored the interplay between meta-learning and in-context learning. The In-context Learning Distillation (ICLD) approach [31] combines in-context learning objectives with language modeling objectives to distill the in-context few-shot learning ability from large pre-trained language models to smaller models. This allows the smaller models to effectively leverage in-context learning while benefiting from the knowledge acquired by the larger models during meta-training.

In addition to few-shot learning and meta-gradient optimization, other meta-learning techniques, such as meta-reinforcement learning [83] and meta-prompt learning [84], have also been explored in the context of in-context learning. These approaches aim to learn meta-strategies or meta-representations that can be efficiently applied to new tasks, enabling the model to rapidly adapt its in-context learning capabilities.

One key challenge in applying meta-learning techniques to in-context learning is the potential for overfitting to the meta-training tasks, which can compromise the model’s ability to generalize to new, unseen tasks. To address this, researchers have explored strategies such as meta-gradient

regularization [84], task-agnostic prompt initialization [85], and diverse meta-training task distributions [81].

Furthermore, the interplay between meta-learning and other learning paradigms, such as transfer learning, few-shot learning, and Bayesian inference, is an active area of research. Understanding the theoretical connections and potential synergies between these approaches can lead to more robust and versatile in-context learning systems [86].

In summary, meta-learning techniques, such as few-shot learning and meta-gradient optimization, have shown great promise in advancing the capabilities of in-context learning. By leveraging prior experience across multiple tasks, these meta-learning approaches can learn effective prompt initialization and adaptation strategies that enable models to rapidly and effectively adapt to new tasks using only a few demonstration examples. As the field of in-context learning continues to evolve, further integration of meta-learning principles and techniques will likely play a crucial role in developing more versatile and adaptable AI systems.

4.4 Addressing Label Biases

In-context learning has emerged as a powerful paradigm that allows large language models (LLMs) to adapt to new tasks using only a few demonstration examples. However, a significant challenge in in-context learning is the sensitivity of model performance to the distribution and characteristics of the demonstration examples, particularly the label biases inherent in these examples.

Recent studies have identified various types of label biases that can arise in in-context learning, including vanilla-label bias, context-label bias, and domain-label bias [87]. These biases can lead to LLMs relying on spurious correlations in the demonstration data rather than truly learning the underlying task, resulting in suboptimal and unreliable performance.

To address this issue, researchers have proposed several techniques to mitigate the sensitivity of in-context learning to label biases in the demonstration examples. One promising approach is the use of calibration methods, which aim to adjust the model’s predictions to account for the biases present in the demonstration data.

[88] introduces a method called Batch Calibration (BC), which controls the contextual bias from the batched input, unifies various prior approaches, and effectively addresses the issue of label biases in in-context learning. BC is a zero-shot, inference-only method that incurs negligible additional computational costs, making it a practical solution for real-world applications.

Another technique, proposed in [87], is domain-context calibration, which estimates the language model’s label bias using random in-domain words from the task corpus and then controls for this bias when making predictions. The authors demonstrate that this method can significantly improve the in-context learning performance of GPT-J and GPT-3 on a wide range of tasks, particularly those with large domain-label bias.

[89] takes a different approach, introducing a method called Generative Calibration. This technique addresses the issue of label shift, where the in-context model shifts the label marginal distribution $p(y)$ while maintaining a good label conditional distribution $p(x|y)$. Generative Calibration adjusts the in-context predictive distribution by estimating the label marginal using Monte-Carlo sampling over the in-context model, effectively mitigating the impact of label biases.

Beyond calibration methods, researchers have also explored bias-aware prompt design as a way to address label biases in in-context learning. [90] introduces a metric to evaluate the predictive bias of a fixed prompt against labels or a given attributes, and then proposes a greedy search-based strategy to identify near-optimal prompts that reduce the impact of label biases.

[91] takes a different approach, constructing “Comparable Demonstrations” (CDs) by minimally editing the text to flip the corresponding labels. The authors find that CDs can significantly

reduce demonstration bias and improve the performance of in-context learning, particularly in out-of-distribution scenarios.

Another line of research focuses on understanding the underlying mechanisms of label biases in in-context learning. [92] introduces novel metrics, such as Label-Correctness Sensitivity and Ground-truth Label Effect Ratio (GLER), to quantify the impact of ground-truth label demonstrations on in-context learning performance. The authors identify key components, such as the verbosity of prompt templates and the language model size, as controlling factors for achieving more noise-resilient in-context learning.

[93] further explores the sensitivity of in-context learning, finding a strong negative correlation between sensitivity and accuracy. The authors propose a few-shot selective prediction method, called SenSel, that abstains from sensitive predictions to improve the reliability of in-context learning.

Overall, the research on addressing label biases in in-context learning highlights the importance of developing robust and calibrated methods that can effectively mitigate the impact of biases in the demonstration data. By combining techniques like calibration, bias-aware prompt design, and the use of carefully constructed demonstration examples, researchers are working towards more reliable and consistent in-context learning capabilities in large language models.

5 Practical Considerations and Future Directions

5.1 Practical Implications of In-Context Learning

The emergence of in-context learning (ICL) in large language models (LLMs) has brought about a fundamental shift in how AI systems can adapt to new tasks and domains. This novel learning paradigm holds immense promise for practical real-world applications, but it also introduces a range of practical considerations and challenges that must be carefully navigated.

One of the key practical implications of ICL is the potential for significant performance gains on a wide variety of tasks, often achieved with only a few demonstration examples. This efficiency is particularly valuable in scenarios where labeled data is scarce or expensive to acquire, such as in specialized domains or for niche applications [8]. The ability to rapidly adapt to new tasks through ICL can enable AI systems to be more responsive to changing user needs and dynamic environments, making them well-suited for real-time decision-making and personalized experiences.

However, the performance advantages of ICL come with trade-offs that must be carefully considered. The sensitivity of ICL to the selection and ordering of in-context examples [63, 71, 94] can introduce challenges in ensuring consistent and reliable performance, particularly in mission-critical applications where predictability and robustness are paramount. Additionally, the strong influence of model priors and the potential for ossification of biases and inconsistencies in LLM predictions, as observed in affective tasks [95], underscores the need for rigorous evaluation and mitigation strategies when deploying ICL-based systems.

Another key practical consideration is the efficiency of ICL. While the ability to adapt to new tasks without fine-tuning offers significant computational advantages, the resource-intensive nature of training LLMs with ICL capabilities can present challenges in terms of scalability and accessibility [96]. Developing strategies to optimize the training process and reduce the computational footprint of ICL could broaden its applicability and make it more accessible to a wider range of organizations and developers.

Closely related to the efficiency challenge is the need for interpretability and transparency in ICL-based systems. The inherent complexity of LLMs and the “black box” nature of their decision-making processes can hinder the understanding and trust of end-users [97]. Advancing interpretability techniques, such as those proposed in the InterpretCC framework [98], can be

crucial for deploying ICL in domains where explainability and accountability are essential, such as healthcare, finance, and education.

Furthermore, the potential for in-context learning to be influenced by biases and inconsistencies in the training data or demonstration examples [90] highlights the need for robust evaluation and monitoring mechanisms. Developing frameworks for assessing the fairness, safety, and ethical implications of ICL-based systems will be crucial for ensuring their responsible deployment in real-world applications.

Finally, the practical implications of ICL extend beyond the immediate performance and efficiency considerations. The emergence of this learning paradigm also raises broader questions about the nature of intelligence, the role of context in learning, and the potential for hybrid approaches that combine ICL with other learning strategies [2, 13, 47, 99, 5]. Exploring these theoretical and conceptual dimensions can inform the development of more versatile and adaptable AI systems that can seamlessly integrate into human-centric workflows and decision-making processes.

5.2 In-Context Learning and Model Architecture

The relationship between the model architecture and in-context learning (ICL) capabilities has been an active area of research in recent years. As the emergence of large language models (LLMs) has brought ICL to the forefront [35], it is crucial to understand how the underlying architectural choices influence a model’s ability to effectively learn and perform tasks in-context.

One of the key components that has been explored in the context of ICL is attention. Transformers, with their multi-headed attention mechanism, have become the dominant architecture for many NLP tasks, including ICL. The self-attention mechanism allows Transformers to dynamically focus on the most relevant parts of the input during the inference process, which is a crucial capability for effective in-context learning [35].

However, the quadratic complexity of the attention operation in Transformers can limit their scalability, especially when dealing with longer input sequences or a large number of in-context examples [100]. This has motivated the exploration of alternative architectures that can maintain the benefits of attention while addressing its computational limitations.

One such alternative is the state-space model (SSM) architecture, exemplified by the Mamba model [101]. SSMs incorporate gating, convolutions, and input-dependent token selection to mitigate the quadratic cost of multi-head attention. While SSMs have shown competitive performance on standard language modeling tasks, their in-context learning capabilities have been less explored compared to Transformers.

A recent study [101] has sought to address this gap by evaluating the ICL performance of SSMs, with a focus on the Mamba model, against Transformer models across various tasks. The results indicate that SSMs can perform comparably to Transformers in standard regression ICL tasks, and even outperform them in tasks like sparse parity learning. However, SSMs fall short in tasks involving non-standard retrieval functionality, highlighting the need for further architectural advancements to address these limitations.

To address the shortcomings of both Transformers and SSMs, researchers have explored hybrid architectures that combine the strengths of different approaches. For example, the MambaFormer model [101] combines the Mamba architecture with attention blocks, surpassing the individual models in tasks where they struggled independently. This suggests that hybrid architectures offer promising avenues for enhancing ICL in language models.

Beyond attention-based architectures, there have been efforts to explore the role of other architectural components in supporting in-context learning. For instance, the importance of positional encoding has been highlighted in the context of linear regression tasks with unstructured

data [102]. The study found that positional encoding can further improve the in-context learning performance of Transformer-based models, indicating that the architectural design choices beyond just attention can have a significant impact on ICL capabilities.

Moreover, the interpretability of the attention mechanism and its role in in-context learning have also been the subject of investigation. A structural analysis of a multi-task question-answering model [103] revealed that attention heads can specialize in particular tasks and that some heads are more conducive to learning than others, both in the multi-task and single-task settings. This suggests that a better understanding of the internal workings of attention-based models can provide valuable insights into the mechanisms underlying their in-context learning abilities.

The interplay between model architecture and in-context learning is a complex and multifaceted topic. While Transformers have emerged as a dominant architecture for ICL, the exploration of alternative architectures, such as SSMS and hybrid models, highlights the potential for further advancements. Ongoing research in this area aims to uncover the specific architectural components and design choices that can enhance the in-context learning capabilities of AI models, ultimately leading to more versatile and adaptable systems that can learn and perform a wide range of tasks with minimal supervision.

5.3 Overcoming Learning Plateaus in In-Context Learning

Overcoming the learning plateaus that often arise during the training of in-context learning capabilities in Transformer models is a critical challenge that has attracted significant attention from the research community. Several recent studies have proposed effective strategies to expedite the learning process and unlock the full potential of in-context learning in these large language models.

One key factor that contributes to the emergence of learning plateaus is the complex interplay between the “weights component” and the “context component” within the model’s internal representations, as highlighted in the paper “Breaking through the learning plateaus of in-context learning in Transformer” [5]. The authors note that the persistence of learning plateaus is often correlated with compromised functionality of the weights component, which they identify as a fundamental driver of this phenomenon. To address this issue, they have developed three strategies that can expedite the learning of in-context learning capabilities in Transformer models: prompt optimization, example selection, and meta-learning techniques.

Prompt optimization, discussed in the subsection “4.1 Prompt Optimization”, is a crucial approach for improving the effectiveness and robustness of in-context learning. By optimizing the prompts used to elicit the desired in-context learning behavior, researchers have demonstrated significant performance gains across a variety of tasks. For example, the paper “Breaking through the learning plateaus of in-context learning in Transformer” [5] shows how prompt search, prompt tuning, and prompt encoding can be leveraged to enhance the in-context learning capabilities of Transformer models.

Another promising approach for overcoming learning plateaus is example selection, as explored in the subsection “4.2 Example Selection”. By carefully selecting the most informative in-context examples, researchers have been able to improve the efficiency and performance of in-context learning. Techniques such as similarity-based retrieval, diversity-based selection, and example-based optimization, as described in the paper “Breaking through the learning plateaus of in-context learning in Transformer” [5], can help Transformer models learn more effectively from a limited set of demonstration examples.

Additionally, the subsection “4.3 Meta-Learning Techniques” highlights the potential of meta-learning approaches for expediting the learning of in-context learning capabilities. By leveraging few-

shot learning and meta-gradient optimization techniques, researchers have shown that Transformer models can learn better prompt initialization and in-context learning strategies, as demonstrated in the paper “Breaking through the learning plateaus of in-context learning in Transformer” [5]. These meta-learning techniques can help overcome the learning plateaus by enabling the models to learn more efficient and adaptable in-context learning strategies.

Furthermore, the paper “The Developmental Landscape of In-Context Learning” [104] reveals that in-context learning in Transformers emerges in discrete developmental stages, suggesting that a more nuanced understanding of this phenomenon may be required to effectively overcome the learning plateaus. The authors introduce methods for detecting the milestones that separate these stages, which can provide valuable insights into the underlying mechanisms and guide the design of more targeted strategies for expediting the learning process.

Addressing the sensitivity of in-context learning to label biases in the demonstration examples is another critical aspect of overcoming learning plateaus, as discussed in the subsection “4.4 Addressing Label Biases”. Techniques such as calibration methods and bias-aware prompt design, as outlined in the paper “Breaking through the learning plateaus of in-context learning in Transformer” [5], can help mitigate the impact of such biases and improve the robustness of in-context learning.

Moreover, the paper “In-context Exploration-Exploitation for Reinforcement Learning” [47] introduces the In-context Exploration-Exploitation (ICEE) algorithm, which aims to optimize the efficiency of in-context policy learning in reinforcement learning tasks. By performing an exploration-exploitation trade-off at inference time within a Transformer model, ICEE can solve Bayesian optimization problems more efficiently than previous in-context learning methods, potentially helping to overcome the learning plateaus in RL domains.

In summary, the research community has made significant progress in addressing the challenge of learning plateaus in in-context learning capabilities of Transformer models. Strategies such as prompt optimization, example selection, meta-learning techniques, and addressing label biases have shown promise in expediting the learning process and unlocking the full potential of in-context learning. As the field continues to evolve, further advancements in these areas, coupled with a deeper understanding of the developmental landscape of in-context learning, are expected to lead to even more robust and efficient in-context learning capabilities in large language models.

5.4 Theoretical Foundations of In-Context Learnability

The theoretical foundations of in-context learning have been an active area of research, with researchers exploring various frameworks to understand the underlying mechanisms and the limits of this learning paradigm. One promising approach is the PAC-based (Probably Approximately Correct) framework, which provides a principled way to analyze the learnability of in-context learning tasks.

The PAC learning framework has been widely used to study the theoretical properties of various machine learning algorithms, and researchers have begun to apply this framework to the study of in-context learning as well. Recent work [13] has proposed a PAC-based framework for understanding the learnability of in-context learning, which aims to provide insights into the task identification versus task learning aspects of this learning paradigm.

In this framework, the authors consider a two-stage learning process: an initial pretraining phase, where the model is trained on a distribution of tasks, and a subsequent in-context learning phase, where the model is presented with a few demonstration examples of a new task and must adapt to perform that task effectively. The key idea is to analyze the sample complexity of the in-context learning phase, which reflects the number of demonstration examples required for the model to learn the new task to a desired level of accuracy.

The authors make several important observations within this framework. First, they show that under mild assumptions, when the pretraining distribution is a mixture of latent tasks (a model often considered for natural language pretraining), these tasks can be efficiently learned via in-context learning, even though the model’s weights are unchanged and the input significantly diverges from the pretraining distribution. This suggests that in-context learning is more about identifying the task than about learning it, a result that aligns with recent empirical findings [105, 106].

The authors also provide finite sample complexity results for the in-context learning phase, which reveal the interplay between the complexity of the task being learned and the informativeness of the demonstration examples. Specifically, they show that the sample complexity of in-context learning scales inversely with the “distance” between the new task and the pretraining distribution, as well as the “coverage” of the demonstration examples in the task space. This highlights the importance of carefully designing the demonstration examples to ensure effective in-context learning.

Furthermore, the authors draw connections between the proposed PAC-based framework and other learning paradigms, such as meta-learning and few-shot learning. They discuss how the insights gained from the in-context learnability framework can inform the development of more general theories of learning with limited data, which is a fundamental challenge in modern machine learning.

Beyond the PAC-based approach, researchers have explored other theoretical frameworks to understand the inner workings of in-context learning. For instance, the “Pelican Soup Framework” [107] introduces the notion of a “common sense knowledge base” and the idea of “meaning association” to explain the generalization capabilities of in-context learning, particularly when confronted with distribution shifts.

Another interesting line of work [38] has examined the emergence of in-context learning from a more mechanistic perspective, focusing on the role of the model architecture and the specific sequence of learning steps required to achieve effective in-context adaptation. These studies have revealed that the abrupt transitions observed in in-context learning can be traced back to the sequential learning of nested logits, which is facilitated by the intrinsic properties of attention-based networks.

Taken together, these theoretical investigations provide valuable insights into the fundamental principles underlying in-context learning and its connections to other learning paradigms. By continuing to explore the theoretical foundations of in-context learning, researchers can shed light on its limitations, uncover new avenues for improvement, and ultimately, foster the development of more versatile and adaptable AI systems.

References

- [1] What Can Transformers Learn In-Context A Case Study of Simple Function Classes. arXiv:2208.01066v3, 2022. <https://arxiv.org/abs/2208.01066v3>
- [2] Concept-aware Data Construction Improves In-context Learning of Language Models. arXiv:2403.09703v1, 2024. <https://arxiv.org/abs/2403.09703v1>
- [3] Concept-aware Training Improves In-context Learning Ability of Language Models. arXiv:2305.13775v1, 2023. <https://arxiv.org/abs/2305.13775v1>
- [4] Unified Demonstration Retriever for In-Context Learning. arXiv:2305.04320v2, 2023. <https://arxiv.org/abs/2305.04320v2>

- [5] Breaking through the learning plateaus of in-context learning in Transformer. arXiv:2309.06054v2, 2023. <https://arxiv.org/abs/2309.06054v2>
- [6] Skill-Based Few-Shot Selection for In-Context Learning. arXiv:2305.14210v2, 2023. <https://arxiv.org/abs/2305.14210v2>
- [7] The Impact of Demonstrations on Multilingual In-Context Learning A Multidimensional Analysis. arXiv:2402.12976v1, 2024. <https://arxiv.org/abs/2402.12976v1>
- [8] Measuring Pointwise $\mathcal{V}$ -Usable Information In-Context-ly. arXiv:2310.12300v2, 2023. <https://arxiv.org/abs/2310.12300v2>
- [9] GistScore Learning Better Representations for In-Context Example Selection with Gist Bottlenecks. arXiv:2311.09606v2, 2023. <https://arxiv.org/abs/2311.09606v2>
- [10] Few-shot Fine-tuning vs. In-context Learning A Fair Comparison and Evaluation. arXiv:2305.16938v2, 2023. <https://arxiv.org/abs/2305.16938v2>
- [11] A Closer Look at In-Context Learning under Distribution Shifts. arXiv:2305.16704v1, 2023. <https://arxiv.org/abs/2305.16704v1>
- [12] Grimoire is All You Need for Enhancing Large Language Models. arXiv:2401.03385v2, 2024. <https://arxiv.org/abs/2401.03385v2>
- [13] The Learnability of In-Context Learning. arXiv:2303.07895v1, 2023. <https://arxiv.org/abs/2303.07895v1>
- [14] PRODIGY Enabling In-context Learning Over Graphs. arXiv:2305.12600v1, 2023. <https://arxiv.org/abs/2305.12600v1>
- [15] A Theory of Emergent In-Context Learning as Implicit Structure Induction. arXiv:2303.07971v1, 2023. <https://arxiv.org/abs/2303.07971v1>
- [16] BERT Pre-training of Deep Bidirectional Transformers for Language Understanding. arXiv:1810.04805v2, 2018. <https://arxiv.org/abs/1810.04805v2>
- [17] Self-training Improves Pre-training for Natural Language Understanding. arXiv:2010.02194v1, 2020. <https://arxiv.org/abs/2010.02194v1>
- [18] Scaling Language Models Methods, Analysis & Insights from Training Gopher. arXiv:2112.11446v2, 2021. <https://arxiv.org/abs/2112.11446v2>
- [19] Emergent Abilities in Reduced-Scale Generative Language Models. arXiv:2404.02204v1, 2024. <https://arxiv.org/abs/2404.02204v1>
- [20] How does Multi-Task Training Affect Transformer In-Context Capabilities Investigations with Function Classes. arXiv:2404.03558v1, 2024. <https://arxiv.org/abs/2404.03558v1>
- [21] Meta-in-context learning in large language models. arXiv:2305.12907v1, 2023. <https://arxiv.org/abs/2305.12907v1>
- [22] In-context Interference in Chat-based Large Language Models. arXiv:2309.12727v1, 2023. <https://arxiv.org/abs/2309.12727v1>

- [23] An Explanation of In-context Learning as Implicit Bayesian Inference. arXiv:2111.02080v6, 2021. <https://arxiv.org/abs/2111.02080v6>
- [24] A Comprehensive Overview of Large Language Models. arXiv:2307.06435v9, 2023. <https://arxiv.org/abs/2307.06435v9>
- [25] Towards More Unified In-context Visual Understanding. arXiv:2312.02520v2, 2023. <https://arxiv.org/abs/2312.02520v2>
- [26] Fine-Tune Language Models as Multi-Modal Differential Equation Solvers. arXiv:2308.05061v4, 2023. <https://arxiv.org/abs/2308.05061v4>
- [27] Instruct Me More! Random Prompting for Visual In-Context Learning. arXiv:2311.03648v1, 2023. <https://arxiv.org/abs/2311.03648v1>
- [28] Understanding and Improving In-Context Learning on Vision-language Models. arXiv:2311.18021v1, 2023. <https://arxiv.org/abs/2311.18021v1>
- [29] Meta-Learned Models of Cognition. arXiv:2304.06729v1, 2023. <https://arxiv.org/abs/2304.06729v1>
- [30] Prior-Knowledge and Attention-based Meta-Learning for Few-Shot Learning. arXiv:1812.04955v5, 2018. <https://arxiv.org/abs/1812.04955v5>
- [31] In-context Learning Distillation Transferring Few-shot Learning Ability of Pre-trained Language Models. arXiv:2212.10670v1, 2022. <https://arxiv.org/abs/2212.10670v1>
- [32] Learning to Balance Bayesian Meta-Learning for Imbalanced and Out-of-distribution Tasks. arXiv:1905.12917v3, 2019. <https://arxiv.org/abs/1905.12917v3>
- [33] Binocular Mutual Learning for Improving Few-shot Classification. arXiv:2108.12104v1, 2021. <https://arxiv.org/abs/2108.12104v1>
- [34] Boosting Meta-Training with Base Class Information for Few-Shot Learning. arXiv:2403.03472v1, 2024. <https://arxiv.org/abs/2403.03472v1>
- [35] Is attention required for ICL Exploring the Relationship Between Model Architecture and In-Context Learning Ability. arXiv:2310.08049v3, 2023. <https://arxiv.org/abs/2310.08049v3>
- [36] Understanding In-Context Learning via Supportive Pretraining Data. arXiv:2306.15091v1, 2023. <https://arxiv.org/abs/2306.15091v1>
- [37] Measuring Inductive Biases of In-Context Learning with Underspecified Demonstrations. arXiv:2305.13299v1, 2023. <https://arxiv.org/abs/2305.13299v1>
- [38] The mechanistic basis of data dependence and abrupt learning in an in-context classification task. arXiv:2312.03002v1, 2023. <https://arxiv.org/abs/2312.03002v1>
- [39] Human Curriculum Effects Emerge with In-Context Learning in Neural Networks. arXiv:2402.08674v1, 2024. <https://arxiv.org/abs/2402.08674v1>
- [40] Pretraining task diversity and the emergence of non-Bayesian in-context learning for regression. arXiv:2306.15063v2, 2023. <https://arxiv.org/abs/2306.15063v2>

- [41] What and How does In-Context Learning Learn Bayesian Model Averaging, Parameterization, and Generalization. arXiv:2305.19420v2, 2023. <https://arxiv.org/abs/2305.19420v2>
- [42] The Physical Systems Behind Optimization Algorithms. arXiv:1612.02803v5, 2016. <https://arxiv.org/abs/1612.02803v5>
- [43] Conflict-Averse Gradient Descent for Multi-task Learning. arXiv:2110.14048v2, 2021. <https://arxiv.org/abs/2110.14048v2>
- [44] Variational Multi-Task Learning with Gumbel-Softmax Priors. arXiv:2111.05323v1, 2021. <https://arxiv.org/abs/2111.05323v1>
- [45] Bayesian Uncertainty for Gradient Aggregation in Multi-Task Learning. arXiv:2402.04005v1, 2024. <https://arxiv.org/abs/2402.04005v1>
- [46] How Many Pretraining Tasks Are Needed for In-Context Learning of Linear Regression. arXiv:2310.08391v2, 2023. <https://arxiv.org/abs/2310.08391v2>
- [47] In-context Exploration-Exploitation for Reinforcement Learning. arXiv:2403.06826v1, 2024. <https://arxiv.org/abs/2403.06826v1>
- [48] Neural Analogical Matching. arXiv:2004.03573v5, 2020. <https://arxiv.org/abs/2004.03573v5>
- [49] A Brain-inspired Computational Model for Human-like Concept Learning. arXiv:2401.06471v1, 2024. <https://arxiv.org/abs/2401.06471v1>
- [50] Towards Multimodal In-Context Learning for Vision & Language Models. arXiv:2403.12736v1, 2024. <https://arxiv.org/abs/2403.12736v1>
- [51] Learning To Retrieve Prompts for In-Context Learning. arXiv:2112.08633v2, 2021. <https://arxiv.org/abs/2112.08633v2>
- [52] Transfer learning to decode brain states reflecting the relationship between cognitive tasks. arXiv:2206.03950v3, 2022. <https://arxiv.org/abs/2206.03950v3>
- [53] Context sequence theory a common explanation for multiple types of learning. arXiv:2208.04707v1, 2022. <https://arxiv.org/abs/2208.04707v1>
- [54] Images Speak in Images A Generalist Painter for In-Context Visual Learning. arXiv:2212.02499v2, 2022. <https://arxiv.org/abs/2212.02499v2>
- [55] IMProv Inpainting-based Multimodal Prompting for Computer Vision Tasks. arXiv:2312.01771v1, 2023. <https://arxiv.org/abs/2312.01771v1>
- [56] Exploring Effective Factors for Improving Visual In-Context Learning. arXiv:2304.04748v1, 2023. <https://arxiv.org/abs/2304.04748v1>
- [57] Exploring Diverse In-Context Configurations for Image Captioning. arXiv:2305.14800v6, 2023. <https://arxiv.org/abs/2305.14800v6>
- [58] MMICL Empowering Vision-language Model with Multi-Modal In-Context Learning. arXiv:2309.07915v3, 2023. <https://arxiv.org/abs/2309.07915v3>

- [59] Lost in Translation When GPT-4V(ision) Can’t See Eye to Eye with Text. A Vision-Language-Consistency Analysis of VLLMs and Beyond. arXiv:2310.12520v1, 2023. <https://arxiv.org/abs/2310.12520v1>
- [60] MMBench Is Your Multi-modal Model an All-around Player. arXiv:2307.06281v3, 2023. <https://arxiv.org/abs/2307.06281v3>
- [61] CODIS Benchmarking Context-Dependent Visual Comprehension for Multimodal Large Language Models. arXiv:2402.13607v2, 2024. <https://arxiv.org/abs/2402.13607v2>
- [62] IGLUE A Benchmark for Transfer Learning across Modalities, Tasks, and Languages. arXiv:2201.11732v2, 2022. <https://arxiv.org/abs/2201.11732v2>
- [63] How are Prompts Different in Terms of Sensitivity. arXiv:2311.07230v1, 2023. <https://arxiv.org/abs/2311.07230v1>
- [64] TEMPERA Test-Time Prompting via Reinforcement Learning. arXiv:2211.11890v1, 2022. <https://arxiv.org/abs/2211.11890v1>
- [65] Instance-wise Prompt Tuning for Pretrained Language Models. arXiv:2206.01958v1, 2022. <https://arxiv.org/abs/2206.01958v1>
- [66] Residual Prompt Tuning Improving Prompt Tuning with Residual Reparameterization. arXiv:2305.03937v1, 2023. <https://arxiv.org/abs/2305.03937v1>
- [67] InfoPrompt Information-Theoretic Soft Prompt Tuning for Natural Language Understanding. arXiv:2306.04933v1, 2023. <https://arxiv.org/abs/2306.04933v1>
- [68] CPL Counterfactual Prompt Learning for Vision and Language Models. arXiv:2210.10362v3, 2022. <https://arxiv.org/abs/2210.10362v3>
- [69] Making Pre-trained Language Models End-to-end Few-shot Learners with Contrastive Prompt Tuning. arXiv:2204.00166v1, 2022. <https://arxiv.org/abs/2204.00166v1>
- [70] Visual-Language Prompt Tuning with Knowledge-guided Context Optimization. arXiv:2303.13283v1, 2023. <https://arxiv.org/abs/2303.13283v1>
- [71] NICE To Optimize In-Context Examples or Not. arXiv:2402.06733v2, 2024. <https://arxiv.org/abs/2402.06733v2>
- [72] Finding Support Examples for In-Context Learning. arXiv:2302.13539v3, 2023. <https://arxiv.org/abs/2302.13539v3>
- [73] Coverage-based Example Selection for In-Context Learning. arXiv:2305.14907v3, 2023. <https://arxiv.org/abs/2305.14907v3>
- [74] RetICL Sequential Retrieval of In-Context Examples with Reinforcement Learning. arXiv:2305.14502v2, 2023. <https://arxiv.org/abs/2305.14502v2>
- [75] Structural Learning of Diverse Ranking. arXiv:1504.04596v2, 2015. <https://arxiv.org/abs/1504.04596v2>
- [76] Data Curation Alone Can Stabilize In-context Learning. arXiv:2212.10378v2, 2022. <https://arxiv.org/abs/2212.10378v2>

- [77] Prompt Optimization via Adversarial In-Context Learning. arXiv:2312.02614v2, 2023. <https://arxiv.org/abs/2312.02614v2>
- [78] Online Meta-Learning. arXiv:1902.08438v4, 2019. <https://arxiv.org/abs/1902.08438v4>
- [79] Multimodal Prototypical Networks for Few-shot Learning. arXiv:2011.08899v1, 2020. <https://arxiv.org/abs/2011.08899v1>
- [80] Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks. arXiv:1703.03400v3, 2017. <https://arxiv.org/abs/1703.03400v3>
- [81] MetaICL Learning to Learn In Context. arXiv:2110.15943v2, 2021. <https://arxiv.org/abs/2110.15943v2>
- [82] Gradient-Regulated Meta-Prompt Learning for Generalizable Vision-Language Models. arXiv:2303.06571v2, 2023. <https://arxiv.org/abs/2303.06571v2>
- [83] Variable-Shot Adaptation for Online Meta-Learning. arXiv:2012.07769v1, 2020. <https://arxiv.org/abs/2012.07769v1>
- [84] Self-supervised Meta-Prompt Learning with Meta-Gradient Regularization for Few-shot Generalization. arXiv:2303.12314v4, 2023. <https://arxiv.org/abs/2303.12314v4>
- [85] Meta-Learning with Context-Agnostic Initialisations. arXiv:2007.14658v2, 2020. <https://arxiv.org/abs/2007.14658v2>
- [86] Stress Testing of Meta-learning Approaches for Few-shot Learning. arXiv:2101.08587v1, 2021. <https://arxiv.org/abs/2101.08587v1>
- [87] Mitigating Label Biases for In-context Learning. arXiv:2305.19148v3, 2023. <https://arxiv.org/abs/2305.19148v3>
- [88] Batch Calibration Rethinking Calibration for In-Context Learning and Prompt Engineering. arXiv:2309.17249v2, 2023. <https://arxiv.org/abs/2309.17249v2>
- [89] Generative Calibration for In-context Learning. arXiv:2310.10266v1, 2023. <https://arxiv.org/abs/2310.10266v1>
- [90] Fairness-guided Few-shot Prompting for Large Language Models. arXiv:2303.13217v3, 2023. <https://arxiv.org/abs/2303.13217v3>
- [91] Comparable Demonstrations are Important in In-Context Learning A Novel Perspective on Demonstration Selection. arXiv:2312.07476v2, 2023. <https://arxiv.org/abs/2312.07476v2>
- [92] Ground-Truth Labels Matter A Deeper Look into Input-Label Demonstrations. arXiv:2205.12685v2, 2022. <https://arxiv.org/abs/2205.12685v2>
- [93] On the Relation between Sensitivity and Accuracy in In-context Learning. arXiv:2209.07661v3, 2022. <https://arxiv.org/abs/2209.07661v3>
- [94] In-context Example Selection with Influences. arXiv:2302.11042v2, 2023. <https://arxiv.org/abs/2302.11042v2>
- [95] The Strong Pull of Prior Knowledge in Large Language Models and Its Impact on Emotion Recognition. arXiv:2403.17125v1, 2024. <https://arxiv.org/abs/2403.17125v1>

- [96] The Mystery of In-Context Learning A Comprehensive Survey on Interpretation and Analysis. arXiv:2311.00237v2, 2023. <https://arxiv.org/abs/2311.00237v2>
- [97] Model-Agnostic Interpretation Framework in Machine Learning A Comparative Study in NBA Sports. arXiv:2401.02630v1, 2024. <https://arxiv.org/abs/2401.02630v1>
- [98] InterpretCC Conditional Computation for Inherently Interpretable Neural Networks. arXiv:2402.02933v1, 2024. <https://arxiv.org/abs/2402.02933v1>
- [99] In-Context Probing Toward Building Robust Classifiers via Probing Large Language Models. arXiv:2305.14171v3, 2023. <https://arxiv.org/abs/2305.14171v3>
- [100] Scaling In-Context Demonstrations with Structured Attention. arXiv:2307.02690v1, 2023. <https://arxiv.org/abs/2307.02690v1>
- [101] Can Mamba Learn How to Learn A Comparative Study on In-Context Learning Tasks. arXiv:2402.04248v2, 2024. <https://arxiv.org/abs/2402.04248v2>
- [102] Benefits of Transformer In-Context Learning in Linear Regression Tasks with Unstructured Data. arXiv:2402.00743v1, 2024. <https://arxiv.org/abs/2402.00743v1>
- [103] Structural analysis of an all-purpose question answering model. arXiv:2104.06045v1, 2021. <https://arxiv.org/abs/2104.06045v1>
- [104] The Developmental Landscape of In-Context Learning. arXiv:2402.02364v1, 2024. <https://arxiv.org/abs/2402.02364v1>
- [105] Prompt-Augmented Linear Probing Scaling beyond the Limit of Few-shot In-Context Learners. arXiv:2212.10873v3, 2022. <https://arxiv.org/abs/2212.10873v3>
- [106] Probing via Prompting. arXiv:2207.01736v1, 2022. <https://arxiv.org/abs/2207.01736v1>
- [107] Understanding In-Context Learning with a Pelican Soup Framework. arXiv:2402.10424v1, 2024. <https://arxiv.org/abs/2402.10424v1>